



PyCon Canada 2017



Using Python to make sense of system traces

*Effici***OS**

jeremie.galarneau@efficios.com 
jgalar 

Presenter



Jérémie Galarneau



EfficiOS Inc.

- Vice President
- <http://www.efficios.com>



Maintainer of

- LTTng-tools
- Babeltrace

Outline

- Quick introduction to tracing
- What is LTTng
- Using Python to work with traces
- Q&A

Isn't tracing just another name for logging?

- Tracers are not completely unlike loggers
- Both save information used to understand the state of an application
 - Trying to cater to developers, admins, and end-users
- In both cases, a careful balance must be achieved between verbosity, performance impact, and usefulness
 - Logging levels
 - Enabling only certain events

Different goals, different tradeoffs

- Tracers focus on low-level events
 - syscalls, scheduling, filesystem events, etc. (1000+ events / second)
 - More events to capture means we must lower the cost / event
 - Binary format
 - Different tracers use different strategies to minimize cost

Different goals, different tradeoffs

- Traces are harder to work with than text log files
 - File size
 - Exploration is difficult
 - Must know the application / kernel to make sense of what was captured
 - Purpose-built tools are needed

LTTng: Linux Trace Toolkit Next Generation

Open source tracing framework for Linux first released in 2005

- Regroups a number of projects
 - LTTng-UST
 - LTTng-modules
 - LTTng-tools
 - LTTng-analyses



LTtng

Safe to use in production environments

- Reliability
 - Out of process trace configuration and collection facilities
- Low intrusiveness
 - Lock-free ring buffer
 - Wait-free RCU techniques and data structures
 - Non-blocking when buffers are full
- Dynamic configuration = no need to restart applications

LTtng

Unify many information sources

- Kernel
- C/C++ applications
- Python standard logging module
- Java (log4j and `java.util.logging`)
- Others
 - `/proc/lttng-logger`

User space tracing

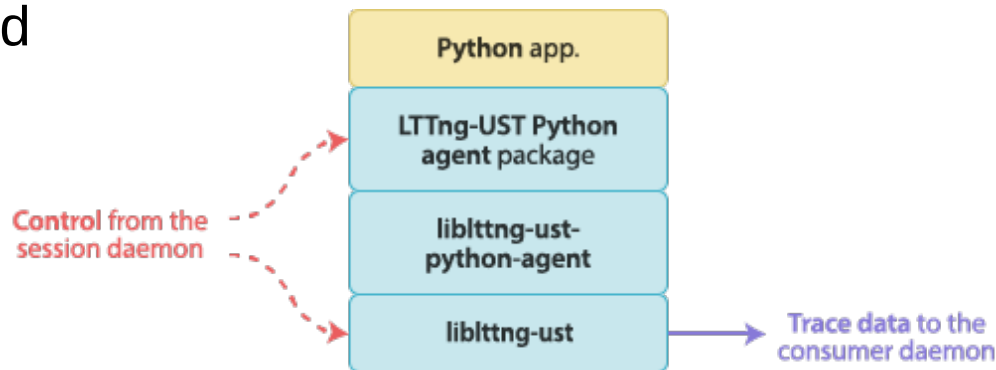
- Useable as a high-performance logging system
 - Can be blocking or non-blocking
- Emitting events from userspace applications provides *anchor points* in kernel traces
- Allows us to figure out *why* an application is performing a certain action

Instrumenting a Python application

LTTng-UST Python agent

- Logging handler for standard **logging** package
- Supports Python 2 and 3

```
import lttngust
```



Better views

- User space developers are building custom tools!
 - Know exactly what they are looking for
 - Modelling their application
 - Tracking internal resources (worker threads, memory pools, connections, users, etc.)
- Originally text-based tools
 - Piping hundreds of GBs of text traces through `grep`, `sed`, `perl`, `awk`...
 - Lots of one-off scripts being passed around
 - Unmanageable, hard to maintain, etc.
 - Break when Babeltrace's text output changes (new event fields)

Babeltrace Python Bindings

- Introduced Python bindings to read traces (2013)
 - Provide users with an easy way to “hack something together”
 - Debugging
 - Testing
 - Reasonably efficient under most scenarios
 - Scripts are maintained as internal tools
- Could we do the same for kernel space?

LTtng analyses

- Development started in early 2014
- Collection of utils
- Models some kernel subsystems to track their current state
 - Latency statistics and distributions (IO, Scheduling, IRQ)
 - System call statistics
 - IRQ handler duration
 - Top resource users

<https://github.com/lttng/lttng-analyses>


Talk about investigating using those scripts: https://www.youtube.com/watch?v=-lUfzp-2_bY

Building our own tools with Python

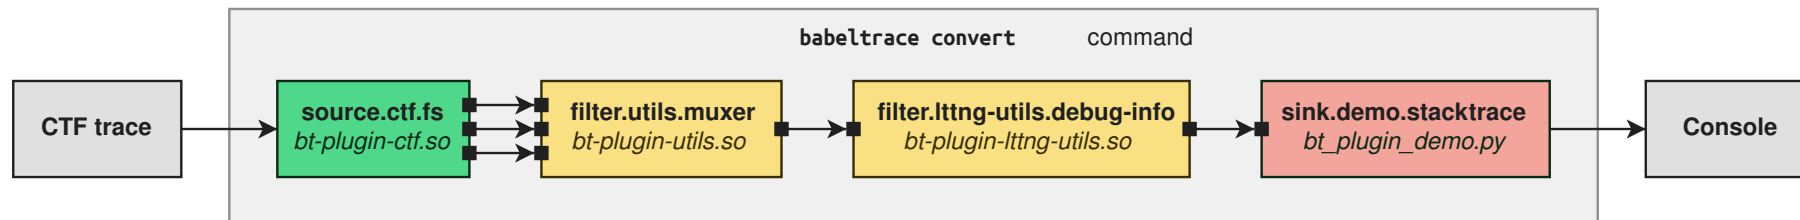
- Python has proven to be a great language to iterate on new tools
- Focus on finding ways to filter through the noise and find clues in traces
- Performance can be a concern if you try to work on too many events
 - Carefully choosing what you trace pays both at runtime and during the analysis

Babeltrace 2.0

- Provides generic components that can be combined to form a trace processing graph
 - CTF file system source
 - CTF file system sink
 - LTTng-live source
 - dmesg source
 - Muxer
 - Trimmer
 - Debugging information injector
- Components can be written in C, C++, and Python



“Callstack” view



Questions ?



🌐 Ittng.org

💬 Ittng-dev@lists.Ittng.org

🐦 [@Ittng_project](https://twitter.com/Ittng_project)

Babeltrace

🌐 diamon.org/babeltrace

🐙 github.com/efficios/babeltrace

💬 Ittng-dev@lists.Ittng.org